



ELSEVIER

Computer Networks 34 (2000) 579–595

**COMPUTER  
NETWORKS**

www.elsevier.com/locate/comnet

# The 1999 DARPA off-line intrusion detection evaluation

Richard Lippmann \*, Joshua W. Haines, David J. Fried, Jonathan Korba,  
Kumar Das

*MIT Lincoln Laboratory, 244 Wood Street, Lexington, MA 02420-9108, USA*

## Abstract

Eight sites participated in the second Defense Advanced Research Projects Agency (DARPA) off-line intrusion detection evaluation in 1999. A test bed generated live background traffic similar to that on a government site containing hundreds of users on thousands of hosts. More than 200 instances of 58 attack types were launched against victim UNIX and Windows NT hosts in three weeks of training data and two weeks of test data. False-alarm rates were low (less than 10 per day). The best detection was provided by network-based systems for old probe and old denial-of-service (DoS) attacks and by host-based systems for Solaris user-to-root (U2R) attacks. The best overall performance would have been provided by a combined system that used both host- and network-based intrusion detection. **Detection accuracy was poor for previously unseen, new, stealthy and Windows NT attacks. Ten of the 58 attack types were completely missed by all systems.** Systems missed attacks because signatures for old attacks did not generalize to new attacks, auditing was not available on all hosts, and protocols and TCP services were not analyzed at all or to the depth required. Promising capabilities were demonstrated by host-based systems, anomaly detection systems and a system that performs forensic analysis on file system data. © 2000 Elsevier Science B.V. All rights reserved.

**Keywords:** Intrusion detection; Evaluate; Attack; Audit; Test bed

## 1. Introduction

The potential damage that can be inflicted by attacks launched over the Internet keeps increasing due to a growing reliance on the Internet and more extensive connectivity. Intrusion detection systems have become an essential component of computer security to detect attacks that occur despite the best preventative measures. Comprehensive discussions of alternate approaches to intrusion detection are available in [1,2,14]. Some

approaches detect attacks in real time and can be used to monitor, and possibly stop, an attack in progress. Others provide after-the-fact forensic information about attacks and can help repair damage, understand the attack mechanism and reduce the possibility of future attacks of the same type. More advanced intrusion detection systems detect never-before-seen, new attacks, while the more typical systems detect previously seen, known attacks.

The widespread deployment and high cost of both commercial and government-developed intrusion detection systems has led to an interest in evaluating these systems. Technical evaluations that focus on algorithm performance are essential for ongoing research. They can contribute to rapid research progress by focusing efforts on difficult

\* Corresponding author. Tel.: +1-617-981-2711; fax: +1-617-981-0186.

E-mail addresses: rpl@sst.ll.mit.edu (R. Lippmann), jhaines@sst.ll.mit.edu (J.W. Haines).

technical areas. They can also produce common, shared corpora or databases which can be used to benchmark performance levels and make it easier for new researchers to enter a field and explore alternate approaches. System evaluations that focus on additional practical issues including cost, ease of use and traffic handling capacity are also useful for determining the capabilities of complete, deployable systems. Without careful evaluations, installing an intrusion detection system could be detrimental because it might lead to a relaxation of vigilance based on unproven assumptions concerning system performance. It might also lead to inefficient use of trained personnel if systems produce many difficult-to-analyze false alarms. A careful assessment of intrusion detection systems is essential to understand their capabilities and limitations and construct an effective security posture that makes use of detection and prevention mechanisms.

It is difficult and costly to perform reliable, systematic evaluations of intrusion detection systems. As a result, few such evaluations have been performed. Table 1 summarizes the characteristics of important evaluations of the past that have compared multiple intrusion detection systems. It includes early studies which describe a methodology that can be used for technical evaluations [4,16,17], the most recent and extensive system evaluation of commercial products that we are aware of [20] and the real-time [5] and off-line [11]<sup>1</sup> components of the 1998 Defense Advanced Research Projects Agency (DARPA) intrusion detection evaluation. The first column in Table 1 provides the first author and date of the study, the second column indicates the number of intrusion detection systems evaluated and the third column provides the number of attack types used and also the number of unique victim machines attacked.

The fourth column indicates whether the study analyzed the number of false alarms produced for normal background traffic and also the duration of background traffic used to measure false-alarm rates. The next column indicates whether stealthy versions of attacks were used in an attempt to evade intrusion detection systems and the final column provides additional comments on the study.

Results are not shown in Table 1 because many studies were informal and did not provide detailed information and because metrics differ widely across studies. The primary performance metric in all studies is the attack detection rate for each attack type used. This metric depends on attack details and the specific version of the intrusion detection system that was tested. It is also insufficient when used alone. It must be combined with false-alarm rates for normal traffic to assess the human workload required to operate intrusion detection systems and dismiss false alarms. False-alarm rates above hundreds per day make a system excessively expensive to deploy, even with high detection accuracy. Unless a system provides extensive forensic information which makes alerts or putative detections easy to analyze, security analysts will not trust alerts and will spend many hours each day dismissing false alarms. Low false-alarm rates combined with high detection rates, however, mean that alerts can be trusted and the human labor required to confirm detections is minimized. **Only recent DARPA evaluations have measured false-alarm rates with a large quantity of rich background traffic.** Other important metrics used by some studies include cost of commercial systems, ease of software installation and use, traffic handling capacity and run-time memory and CPU requirements.

As can be seen from Table 1, evaluations have become more complex and extensive over the years. Initial evaluations included few systems, few attack types, did not include stealthy attacks and included little normal background traffic to evaluate false-alarm rates. The 1998 off-line DARPA evaluation includes 10 systems, 38 attack types, weeks of rich background traffic and stealthy attacks and also led to a corpus or database of attacks and background traffic that is being widely

<sup>1</sup> A public web site of MIT Lincoln Laboratory, <http://www.ll.mit.edu/IST/ideval/index.html>, contains limited information on the 1998 and 1999 evaluations. Follow instructions on this web site or send email to the authors (rpl@sst.ll.mit.edu or jhaines@sst.ll.mit.edu) to obtain access to a password-protected site with more complete information on these evaluations and results. Software scripts to execute attacks are not provided on these or other web sites.

Table 1  
Characteristics of past intrusion detection evaluations

Study	IDs	Attacks/victims	False alarms	Stealth	Comments
Puketza, 1994 [16,17]	2	4/1	Yes/Unknown	No	Automated attacks and simple telnet traffic
Debar, 1998 [4]	3	4/1	Yes/Unknown	No	Automated attacks and FTP traffic
Shipley, 1999 [20]	10	12/4	No/None	Yes	Product comparison of 10 commercial IDs
Durst, 1999 [5]	4	19/4	Yes/Hours	Yes	1998 DARPA real-time evaluation
Lippmann, 2000 [11]	10	38/4	Yes/Weeks	Yes	1998 DARPA off-line evaluation, distribute standard ID corpus

used for the evaluation and development of intrusion detection systems. The first two evaluations in Table 1 describe initial research programs designed to develop a methodology for intrusion detection evaluation [4,16,17]. Both studies incorporated scripting software to provide repeatability by automating generation of attacks and background traffic. Few attack types were used in these studies and background traffic consisted of a small number of automated telnet or FTP sessions. Both studies demonstrated the importance of repeatability for intrusion detection system development. Initial low detection and high false-alarm rates were improved by cyclical testing and development with repeatable attacks and background traffic. The second study [4] also noted that generating realistic normal background traffic was complex and time-consuming in heterogeneous computing environments.

Many product comparisons of commercial intrusion detection systems have been published in the past few years. The third entry in Table 1 is a recent, comprehensive, product evaluation. It included three host-based and seven network-based commercial intrusion detection systems which were evaluated using more than 12 attack types and four victim machines. This study also included stealthy probe or scan attacks and stealthy packet modifications described in [15] designed to elude intrusion detection systems. This study did not provide detailed per-attack detection results, but mentioned that no system detected all attacks and that stealthy attacks successfully eluded many systems. Most of the evaluated systems relied on attack “signatures” to detect old or known attacks. New signatures can often be added by hand or downloaded from a remote site. This evaluation focused on practical system characteristics such as

ease of use and cost and did not measure false-alarm rates for normal background traffic. It did, however, use network load-generating software to demonstrate that some network-based intrusion detection systems fail to detect attacks at high network loads.

The last two rows in Table 1 are for real-time and off-line DARPA 1998 evaluations. As can be seen from the table, the off-line evaluation is the most complex performed to date. It was an initial attempt at a comprehensive evaluation which included background traffic to measure false-alarm rates, many attacks, and more than eight different intrusion detection systems. This exploratory evaluation was limited. It included only intrusion detection systems developed under DARPA sponsorship, attacks against UNIX hosts and background traffic designed to be similar to traffic on one air force base. Six research groups participated in this statistically blind evaluation to provide unbiased measurement of current performance levels. The off-line evaluation, performed by MIT Lincoln Laboratory, included weeks of training and test traffic and more than 300 instances of 38 attack types and resulted in an archival 1998 intrusion detection corpus or database [11].<sup>1</sup> This corpus can be processed simultaneously at many sites to evaluate and develop research systems and it continues to be used for algorithm development and as a baseline for future evaluations. The real-time evaluation, performed by the Air Force Research Laboratory (AFRL), evaluated a smaller number of systems, which have real-time implementations using a more complex network, fewer attacks and four hours of traffic [5]. Results of the 1998 evaluation helped determine the strengths and weaknesses of alternative technical approaches and had a strong influence on

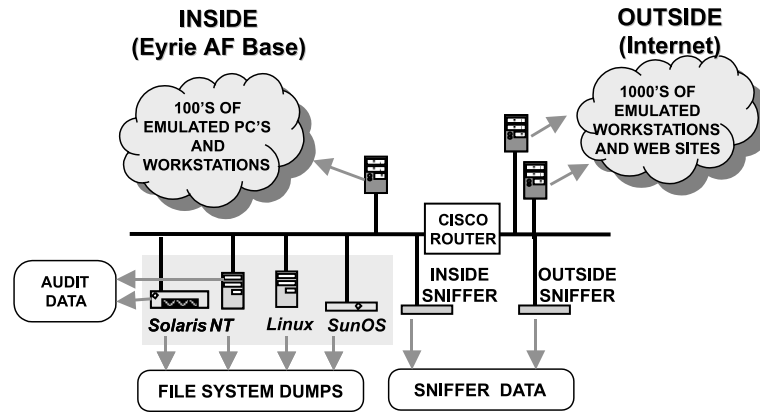


Fig. 1. Block diagram of 1999 test bed.

**DARPA intrusion detection research goals.** Further off-line and real-time evaluations which built on the initial 1998 effort were performed in 1999. This paper reports the results of the off-line 1999 evaluation. Results and lessons learned from the 1998 off-line evaluation are first summarized, the 1999 off-line evaluation is described, 1999 **results** are presented and suggestions are provided for future evaluations. Further details of the 1999 off-line evaluation are available in [3,10,12].<sup>1</sup>

## 2. Summary of the 1998 off-line evaluation

The DARPA 1998 Intrusion Detection Evaluation was an initial attempt to perform a comprehensive technical evaluation of intrusion detection technology. As noted above, this evaluation had limited goals. It was designed to evaluate only DARPA-funded intrusion detection technology and not complete deployable intrusion detection systems or commercial systems. It was also designed to measure false-alarm rates using background traffic similar to that on one air force base and detection rates of remotely-initiated attacks against UNIX hosts. Fig. 1 shows the current version of an isolated test bed network which was first developed for the 1998 off-line evaluation. Scripting techniques which extend the approaches used in [4,16] are used to generate live background traffic which is similar to traffic that flows between the inside of an air force base and the outside Internet.

This approach was selected for the evaluation because hosts can be attacked without degrading operational Air Force systems, and corpora containing background traffic and attacks can be widely distributed without security or privacy concerns. A rich variety of background traffic, which looks as if it were initiated by hundreds of users on thousands of hosts, is generated in the test bed. The left side of Fig. 1 represents the inside of the fictional Eyrie Air Force Base created for the evaluations and the right side represents the outside Internet. The 1998 evaluation did not include the Windows NT victim machine or the inside sniffer shown on the left of Fig. 1, but instead focused exclusively on UNIX and router attacks. Automated attacks were launched against three inside UNIX victim machines (SunOS, Solaris, Linux) and the router from outside hosts. More than 300 instances of 38 different attacks were embedded in seven weeks of training data and two weeks of test data. Machines labeled “sniffer” in Fig. 1 ran a program named *tcpdump*<sup>2</sup> to capture all packets transmitted over the attached network segment.

Six research sites participated in the blind 1998 evaluation and the results were analyzed to determine attack detection rate as a function of

<sup>2</sup> The Lawrence Berkeley National Laboratory Network Research Group provides *tcpdump* at <http://www-nrg.ee.lbl.gov/>

false-alarm rate. Performance was evaluated for old attacks included in the training data and new attacks which only occurred in the test data. **Detection performance for the best systems was more than 60% correct at and below a false-alarm rate of 10 false alarms per day for both old and new probe attacks and attacks where a local user illegally became root (U2R).** Detection rates were mixed for denial-of-service (DoS) attacks and remote-to-local (R2L) attacks where a remote user illegally accessed a local host. Although detection accuracy for old attacks in these two categories was roughly 80%, detection accuracy for new and novel attacks was below 25% even at high false-alarm rates. **These results demonstrated that current intrusion detection systems do not detect new attacks well and refocused research goals on techniques which can detect new attacks.** Results of the real-time evaluation generally agreed with those of the off-line evaluation. Detection rates for the systems and attacks in common were similar. Two interesting results from the off-line evaluation were that slow, stealthy scans were not well detected by some intrusion detection systems and false-alarm rates of a network-based system used by the Air Force were similar to those of a reference keyword-based system used in the off-line evaluation.

### 3. Conclusions from the 1998 evaluation

The 1998 evaluation uncovered a widespread interest in obtaining training and test corpora containing normal traffic and attacks to develop and evaluate intrusion detection systems. To date, more than 90 sites have downloaded all or part of the 1998 off-line intrusion detection archival corpus from a Lincoln Laboratory web site.<sup>1</sup> Information from sites which have downloaded this corpus indicates that it is being used to evaluate and develop both commercial and research intrusion detection systems (e.g., [21]) and to train security analysts. A processed subset of this corpus was also redistributed as part of a contest sponsored by the International Conference on Knowledge Discovery in Databases [6]. This conference attracted 24 participants, who used modern approaches to pattern classification to achieve high

performance on a constrained intrusion detection task.

The 1998 evaluation also demonstrated that it is possible to evaluate a diverse collection of intrusion detection systems but that this is more complex than initial analyses suggested. All components of the evaluation, from designing and managing the test bed to generating background traffic, scoring systems, automating, running, marking ground truth and verifying attacks, included added complexities caused by the wide variety of traffic, attacks and intrusion detection systems included. For example, labeling attacks involved annotating every network packet associated with each attack. This was partially automated, but it required extensive hand correction and analyses which had to be customized for each attack. Experiences of the Lincoln Laboratory evaluators led to suggestions for reducing the cost and complexity of the evaluation. These included simplifying scoring procedures, requesting more detailed and formal system descriptions from participants, automating attack generation and verification more fully and automating more of the daily procedures required to run the test bed continuously. Experiences by the many participants and others also led to suggestions for improving the evaluation process. These included providing training data containing no attacks to train anomaly detectors, simplifying scoring procedures, exploring false-alarm rates with a richer range of background traffic, providing a written security policy and performing more detailed analyses of misses and false alarms. All of these suggestions were incorporated in the 1999 evaluation.

### 4. Overview of the 1999 evaluation

The 1999 evaluation was a blind off-line evaluation, as in 1998, but modified based on suggestions from 1998 and also with major extensions to enhance the analysis and cover more attack types. Fig. 1 shows a block diagram of the 1999 test bed. Major changes for 1999 are the addition of a Windows NT workstation as a victim, the addition of an inside tcpdump sniffer machine and the collection of both Windows NT audit events and

inside tcpdump sniffing data for inclusion in archival data provided to participants. Not shown in this figure are new Windows NT workstations added to support NT attacks, new inside attacks and new stealthy attacks designed to avoid detection by network-based systems tested in 1998. The Windows NT victim machine and the associated attacks and audit data were added due to increased reliance on Windows NT systems by the military. Inside attacks and inside sniffer data to detect these attacks were added due to the dangers posed by inside attacks. Stealthy attacks were added due to an emphasis on sophisticated attackers who can carefully craft attacks to look like normal traffic. In addition, two new types of analyses were performed. First, an analysis of misses and high-scoring false alarms was performed for each system to determine why systems miss specific attacks and what causes false alarms. Second, participants were optionally permitted to submit attack forensic information that could help a security analyst identify important characteristics of the attack and respond. This identification information included the attack category, the names of old attacks, the ports/protocols used and the IP addresses used by the attacker.

Another major change in 1999 was a focus on determining the ability of systems to detect new attacks without first training on instances of these attacks. The 1998 evaluation demonstrated that systems could not detect new attacks well. The new 1999 evaluation was designed to evaluate enhanced systems which could detect new attacks and to analyze why systems missed new attacks. Many new attacks were thus developed and only examples of a few of these were provided in training data.

## 5. Test bed network and background traffic

The inside of the simulated Eryie Air Force Base shown in Fig. 1 contains four primary victim machines (Linux 2.0.27, SunOS 4.1.4, Sun Solaris 2.5.1, Windows NT 4.0), a sniffer to capture network traffic and a gateway to hundreds of other inside emulated PCs and workstations. The outside simulated Internet contains a sniffer, a gate-

way to hundreds of emulated workstations on many other subnets and a second gateway to thousands of emulated web servers. Data collected to evaluate intrusion detection systems include network sniffing data from the inside and outside sniffers, Solaris Basic Security Module (BSM) audit data collected from the Solaris host, Windows NT audit event logs collected from the Windows NT host, nightly listings of all files on the four victim machines and nightly dumps of security-related files on all victim machines.

Custom software automata in the test bed simulate hundreds of programmers, secretaries, managers and other types of users running common UNIX and Windows NT application programs. In addition, custom Linux kernel modifications provided by the AFRL allow a small number of actual hosts to appear as if they are thousands of hosts with different IP addresses. Fig. 2 shows the average number of connections per day for the most common TCP services. As can be seen, Web traffic dominates, but many other types of traffic are generated which use a variety of services. User automata send and receive mail, browse web sites, send and receive files using FTP, use telnet and ssh to log into remote computers and perform work, monitor the router remotely using SNMP and perform other tasks. In addition to automatic traffic, the test bed allows human actors to generate background traffic and attacks when the traffic or attack is too complex to automate. Background traffic characteristics, including the overall traffic level, the proportion of traffic from different services and the variability of




Fig. 2. Average connections per day for dominant TCP services.



traffic with time of day, are similar to characteristics measured on a small air force base in 1998. The average number of background-traffic bytes transmitted per day between the inside and outside of this test bed is roughly 411 Mb per day, with most of the traffic concentrated between 8:00 AM and 6:00 PM. The dominant protocols are TCP (384 Mb), UDP (26 Mb) and ICMP (98 kb). These traffic rates are low compared to current rates at some large commercial and academic sites, but are representative of traffic measured at the beginning of this project. These rates also lead to sniffed data file sizes that can still be transported over the Internet without practical difficulties. **The flat test bed structure without firewalls or other protective devices simplifies maintenance and attack generation.** Future evaluations will include firewalls, more complex architectures, attacks against firewalls and more complex attacks including man-in-the-middle attacks that take advantage of a network hierarchy.

## 6. Attacks

Twelve new Windows NT attacks were added in 1999 along with stealthy versions of many 1998 attacks, new inside console-based attacks and six

new UNIX attacks. The 56 different attack types shown in Tables 2 and 3 were used in the evaluation. Attacks in normal font in these tables are old attacks from 1998 executed in the clear (114 instances). Attacks in italics are new attacks developed for 1999 (62 instances) or stealthy versions of attacks used in 1998 (35 instances). Details of attacks, including further references and information on implementations, are available in [3,9,10,12]. Five major attack categories and the attack victims are shown in Tables 2 and 3. The primary victims, listed along the top of these tables, are the four inside victim hosts, shown in the gray box of Fig. 1, and the Cisco router. In addition, some probes query all machines in a given range of IP addresses as indicated by the column labeled “all” in Table 2.

The upper row of Table 2 lists probe or scan attacks. These attacks automatically scan a network of computers or a DNS server to find valid IP addresses (ipsweep, lsdomain, mscan), active ports (portsweep, mscan), host operating system types (queso, mscan) and known vulnerabilities (satan). All of these probes, except two (mscan and satan), are either new in 1999 (e.g., ntinfoscan, queso, illegalsniffer) or are stealthy versions of 1999 probes (e.g., portsweep, ipsweep). Probes are considered stealthy if they issue 10 or fewer

Table 2  
Probe and denial of service (DoS) attacks

	Solaris	SunOS	NT	Linux	All
<b>Probe (37)</b>	<i>portsweep</i> <i>queso</i>	<i>portsweep</i> <i>queso</i>	<i>ntinfoscan</i> <i>portsweep</i>	<i>lsdomain</i> <i>mscan</i> <i>portsweep</i> <i>queso</i> <i>satan</i>	<i>illegal-sniffer</i> <i>ipsweep</i> <i>portsweep</i>
<b>DoS (65)</b>	<i>neptune</i> <i>pod</i> <i>processtable</i> <i>selfping</i> <i>smurf</i> <i>syslogd</i> <i>tcpreset</i> <i>warezclient</i>	<i>arppoisson</i> <i>land</i> <i>mailbomb</i> <i>neptune</i> <i>pod</i> <i>processtable</i>	<i>arppoisson</i> <i>crashiis</i> <i>dosmuke</i> <i>smurf</i> <i>tcpreset</i>	<i>apache2</i> <i>arppoisson</i> <i>back</i> <i>mailbomb</i> <i>neptune</i> <i>pod</i> <i>processtable</i> <i>smurf</i> <i>tcpreset</i> <i>teardrop</i> <i>udpstorm</i>	

Table 3

Remote to Local (R2L), User to Root (U2R), and Data attacks

	Solaris	SunOS	NT	Linux	Cisco
R2L (56)	dict ftppwrite guest httptunnel xlock xsnoop	dict xsnoop	dict <i>framespoof</i> <i>netbus</i> <i>netcat</i> <i>ppmacro</i>	dict imap named <i>ncftp</i> phf sendmail <i>sshtrojan</i> xlock xsnoop	snmpget
U2R (37)	<i>eject</i> <i>fdformat</i> <i>ffbconfig</i> <i>ps</i>	<i>loadmodule</i>	<i>casesen</i> <i>ntfsdos</i> <i>nukepw</i> <i>sechole</i> <i>yaga</i>	<i>perl</i> <i>sqlattack</i> xterm	
Data (13)	<i>secret</i>		<i>ntfsdos</i> <i>ppmacro</i>	<i>secret</i> sqlattack	

connections or packets or if they wait longer than 59 s between successive network transmissions. **Stealthy probes are similar to clear probes because they gather similar information concerning IP addresses, vulnerable ports and operating system types. They differ because this information is gathered at a slower rate and because less, but more focused, information is gathered from each attack instance.** For example, stealthy port sweeps are slow and focus only on ports with known vulnerabilities. The new “illegalsniffer” attack is different from the other probes. During this attack, a Linux sniffer machine is installed on the inside network running the tcpdump program in a manner that creates many DNS queries from this new and illegal IP address.

The second row of Table 2 contains DoS attacks designed to disrupt a host or network service. New 1999 DoS attacks crash the Solaris operating system (selfping), actively terminate all TCP connections to a specific host (tcpreset), corrupt ARP cache entries for a victim not in others’ caches (arppoisson), crash the Microsoft Windows NT web server (crashiis) and crash Windows NT (dosnuke).

The first row of Table 3 contains R2L attacks. In these attacks, an attacker who does not have an account on a victim machine gains local access to

the machine (e.g., guest, dict), exfiltrates files from the machine (e.g., ppmacro) or modifies data in transit to the machine (e.g., framespoof). New 1999 R2L attacks include an NT PowerPoint macro attack (ppmacro), a man-in-the middle web browser attack (framespoof), an NT trojan-installed remote-administration tool (netbus), a Linux trojan SSH server (sshtrojan) and a version of a Linux FTP file access-utility with a bug that allows remote commands to run on a local machine (ncftp). The second row of Table 3 contains U2R attacks where a local user on a machine is able to obtain privileges normally reserved for the UNIX super user or the Windows NT administrator. All five NT U2R attacks are new this year and all other attacks except one (xterm) are versions of 1998 U2R attacks that were redesigned to be stealthy to network-based intrusion detection systems evaluated in 1998. Techniques used to make these U2R attacks stealthy are described in [3,10,12]. They include running the attack over multiple sessions, embedding the attack in normal user actions, writing custom buffer-overflow machine code that does not spawn a root-level shell but simply “chmod’s” a file, bundling the complete attack into one shell script, setting up delayed “time bomb” attacks and transferring the attack and the attack output using common network



services. The bottom row in Table 3 contains Data attacks. This is a new attack type added in 1999. The goal of a Data attack is to exfiltrate special files which the security policy specifies should remain on the victim hosts. These include “secret” attacks, where a user who is allowed to access the special files exfiltrates them via common applications such as mail or FTP, and other attacks where privilege to access the special files is obtained using a U2R attack (ntfsdos, sqlattack). Note that an attack could be labeled as both U2R and Data if one of the U2R attacks was used to obtain access to the special files. The “Data” category thus specifies the goal of an attack rather than the attack mechanism.

Attack implementation was simplified for U2R attacks in 1999 by integrating attack automation software with the automaton used to generate telnet sessions. This made it easier to embed attacks within normal telnet sessions. In addition, attack verification was simplified by running all attacks from a separate dedicated machine and sniffing traffic to and from that machine. This made it easier to collect the network traffic generated by each attack. Custom software was required to change routing tables in the test bed gateways whenever the IP address of the dedicated attacker machine changed. This made it possible to isolate network traffic generated by attacks for all but inside attacks which were launched from the console of a victim and for attacks which installed trojans or other types of malicious software on inside machines. Any network traffic for these two types of attacks had to be extracted from inside sniffer data by hand.

## 7. Participants and scoring

Eight research groups participated in the evaluation using a variety of approaches to intrusion detection. Papers by these groups describing high-performing systems are provided in [7,8,13,18, 19,22–24]. One requirement for participation in the evaluation was the submission of a detailed system description, which was used for scoring and analysis. System descriptions described the types of attacks the system was designed to detect, data

sources used, features extracted and whether optional attack identification information was provided as an output. Most systems used network sniffer data to detect Probe and DoS attacks against all systems [8,13,19,23] or BSM Solaris host audit data to detect Solaris R2L and U2R attacks [7,13,23]. Two systems produced a combined output from both network sniffer data and host audit data [13,23]. A few systems used network sniffer data to detect R2L and U2R attacks against the UNIX victims [13,23]. One system used NT audit data to detect U2R and R2L attacks against the Windows NT victim [18] and two systems used BSM audit data to detect Data attacks against the Solaris victim [13,23]. A final system used information from a nightly file system scan to detect R2L, U2R and Data attacks against the Solaris victim [22]. The software program that performs this scan was the only custom auditing tool used in the evaluation. A variety of approaches was employed, including expert systems that use rules or signatures to detect attacks, anomaly detectors, pattern classifiers, recurrent neural networks, data mining techniques and a reasoning system that performs a forensic analysis of the Solaris file system.

Three weeks of training data, comprising two weeks of background traffic with no attacks and one week of background traffic with a few attacks, were provided to participants from mid-May to mid-July 1999 to support system tuning and training. The locations of attacks in the training data were clearly labeled. Two weeks of unlabeled test data were provided from late September to the middle of October. Participants downloaded these data from a web site, processed it through their intrusion detection systems and generated putative hits or alerts at the output of their intrusion detection systems. Lists of alerts were due by early October. In addition, participants could optionally return more extensive identification lists for each attack.

A simplified approach was used in 1999 to label attacks and score alerts and new scoring procedures were added to analyze the optional identification lists. In 1998, essentially every network TCP/IP connection, UDP packet and ICMP packet was labeled, and participants determined

which connections and packets corresponded to attacks. Although this approach pre-specifies all potential attack packets and thus simplifies scoring and analysis, it can make submitting alerts difficult because aligning alerts with the network connections and packets that generate alerts is often complex. In addition, this approach cannot be used with inside attacks that generate no network traffic. In 1999, a new simplified approach was adopted. Each alert only had to indicate the date, time, victim IP address and score for each putative attack detection. An alert could also optionally indicate the attack category. This was used to assign false alarms to attack categories. Putative detections returned by participants were counted as true “hits” or true detections if the time of any alert occurred during the time of any attack segment and the alert was for the correct victim IP address. Alerts that occurred outside all attack segments were counted as “misses” or false alarms. Attack segments correspond to the duration of all network packets and connections generated by an attack and to time intervals when attack processes are running on a victim host. To account for small timing inconsistencies across hosts, extra leeway of 60 s was typically allowed for alerts before and after the end of each attack segment. The analysis of each system only included attacks which that system was designed to detect, as specified in the system description. Systems were not penalized for missing attacks they were not designed to detect and false alarms that occurred during segments of out-of-spec attacks were ignored.

The score produced by a system was required to be a number that increased as the certainty of an attack at the specified time increased. All participants returned numbers ranging between zero and one, and many participants produced binary outputs (0’s and 1’s only). If alerts occurred in multiple attack segments of one attack, then the score assigned to that attack for further analysis was the highest score in all the alerts. Some participants returned optional identification information for attacks. This included attack category, names of old attacks selected from a list of provided names, attack source and destination IP addresses, start time, duration and ports/services used. This

information was analyzed separately from the alert lists used for detection scoring. Results in this paper focus on detection results derived from the required alert lists.

Attack labels were used to designate attack segments in the training data and also to score lists of alerts returned by participants. Attack labels were provided using list files similar to those used in 1998. In addition, a separate list file was provided for each attack specifying all segments of that attack. Entries in these list files included date, start time, duration, unique attack identifier, attack name, source and destination ports and IP addresses, protocol and details concerning the attack. Details included indications that the attack was clear or stealthy, old or new and inside or outside, victim machine type and whether traces of the attack occurred in each of the different data types that was collected.

## 8. Results

An initial analysis was performed to determine how well all systems taken together detect attacks regardless of false-alarm rates. The best system was first selected for each attack as the system which detected the most instances of that attack. The detection rates of these best systems provided a rough upper bound on composite system performance. Thirty seven of the 58 attack types were detected well by this composite system, but many stealthy and new attacks were always or frequently missed. Poorly detected attacks for which half or more of the attack instances were not detected by the best system are listed in Table 4. This table lists the attack name, the attack category, details concerning whether the attack is old, new or stealthy, the total number of instances of this attack and the number of instances detected by the system which detected this attack best. Table 4 contains 21 attack types and is dominated by new attacks and attacks designed to be stealthy to 1998 network-based intrusion detection systems. All instances of 10 of the attack types in Table 4 were totally missed by all systems. These results suggest that the new systems developed for the 1999 evaluation still do

Table 4

Poorly detected attacks where the best system for each attack detects half or fewer of the attack instances

Name	Category	Details	Total instances	Instances detected by best system
ipsweep	Probe	Stealthy	3	0
lsdomain	Probe	Stealthy	2	1
portsweep	Probe	Stealthy	11	3
queso	Probe	New	4	0
resetscan	Probe	Stealthy	1	0
arppoisn	DoS	New	5	1
dosnuke	DoS	New–NT	4	2
selfping	DoS	New	3	0
tcpreset	DoS	New	3	1
warezclient	DoS	Old	3	0
ncftp	R2L	New	5	0
netbus	R2L	New–NT	3	1
netcat	R2L	New–NT	4	2
snmpget	R2L	Old	4	0
sshtrojan	R2L	New	3	0
loadmodule	U2R	Stealthy	3	1
ntfsdos	U2R	New–NT	3	1
perl	U2R	Stealthy	4	0
sechole	U2R	New–NT	3	1
sqlattack	U2R	Stealthy	3	0
xterm	U2R	Old	3	1

not detect new attacks well and that stealthy probes and U2R attacks can avoid detection by network-based systems.

Further analyses evaluated system performance at false-alarm rates in a specified range. The detection rate of each system at different false-alarm rates can be determined by lowering a threshold from 1.0 to 0.0, counting the detections with scores above the threshold as hits and counting the number of alerts above the threshold that do not detect attacks as false alarms. This results in one or more operating points for each system which trade off false-alarm rate against detection rate. It was found that almost all systems, except some anomaly detection systems, achieved their maximum detection accuracy at or below 10 false alarms per day on the 1999 corpus. These low false-alarm rates were presumably due to the low overall traffic volume, the relative stationarity of the traffic and the ability to tune systems on three weeks of training data to reduce false alarms. In the remaining presentation, the detection rate reported for each system is the highest detection rate achieved at or below 10 false alarms per day on the two weeks of test data.

Table 5 shows average detection rates at 10 false alarms per day for each attack category and victim type. This table provides overall results and does not separately analyze old, new and stealthy attacks. The upper number in a cell, surrounded by dashes, is the number of attack instances in that cell and the other entries provide the percent correct detections for all systems with detection rates above 40% in that cell. A cell contains only the number of instances if no system detected more than 40% of the instances. Only one entry is filled for the bottom row because only probe attacks were against all the victim machines and the SunOS/Data cell is empty because there were no Data attacks against the SunOS victim. High-performance systems listed in Table 5 include rule-based expert systems that use network sniffing data and/or Solaris BSM audit data (Expert-1 through Expert-3 [13,19,23]), a data mining system that uses network sniffing data (DMine [8]), a pattern classification approach that uses network sniffing data (PClassify), an anomaly detection system which uses recurrent neural networks to analyze system call sequences in Solaris BSM audit data (Anomaly [7]) and a reasoning system which performs a

Table 5

Percent attack instances detected for systems with a detection rate above 40% in each cell and at false alarm rates below 10 false alarms per day

	DoS	Probe	R2L	U2R	Data
Solaris	-19- Expert-1: 63% Expert-2: 53%	-5- Expert-2: 60% Expert-3: 50%	-12- Expert-1: 50% Forensics: 50%	-11- Expert-1: 100% Expert-2: 100% Anomaly: 100% Forensics: 73%	-6- Expert-2: 100% Forensics: 83%
NT	-16- Expert-1: 69% Expert-2: 69%	-5- Expert-1: 80% Expert-2: 60%	-12-	-13-	-5-
SunOS	-8- DMine: 88% Expert-1: 63% Expert-2: 50%	-5- PClassify: 60%	-3- Expert-2: 67%	-3-	
Linux	-19- Expert-1: 84% DMine: 74% Expert-2: 68%	-8- Expert-3: 60% DMine: 50%	-25- Expert-2: 64% Expert-1: 44%	-10-	-4-
All		-11- Expert-1: 46%			

nightly forensic analysis of the Solaris file system (Forensics [22]).

No one approach or system provides the best performance across all categories. The best performance is provided for probe and DoS attacks by systems that use network sniffer data and for U2R and Data attacks against the Solaris victim by systems that use BSM audit data. Detection rates for U2R and Data attacks are generally poor for SunOS and Linux victims where extensive au-

dit data are not available. Detection rates for R2L, U2R and Data attacks are poor for Windows NT, which was included in the evaluation for the first time this year.

Fig. 3 shows the performance of the best intrusion detection system in each attack category at a false-alarm rate of 10 false alarms per day. The left chart compares the percentage of attack instances detected for old-clear and new attacks and the right chart compares performance for

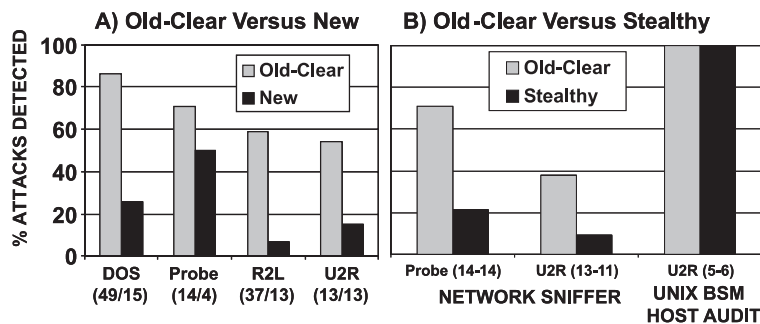


Fig. 3. Comparison of detection at 10 false alarms per day for (A) old-clear versus new attacks and (B) old-clear versus stealthy attacks.

old-clear and stealthy attacks. The numbers in parentheses on the horizontal axis below the attack category indicate the number of instances of attacks of different types. For example, in Fig. 3(A), there were 49 old-clear and 15 new DoS attacks. Fig. 3(A) demonstrates that detection of new attacks was much worse than detection of old-clear attacks across all attack categories, especially for DoS, R2L and U2R attacks. The average detection rate for old-clear attacks was 72% and this dropped to 19% for new attacks. Fig. 3(B) demonstrates that stealthy probes and U2R attacks were much more difficult to detect for network-based intrusion detection systems that used sniffing data. Those attacks against the Solaris victim, however, were accurately detected by host-based intrusion detection systems that used BSM audit data.

Attacks were detected best when they produced a consistent “signature” or sequence of events in tcpdump data or in audit data that was different from sequences produced by normal traffic. A detailed analysis by participants demonstrated that attacks were missed for a variety of reasons. Systems which relied on rules or signatures missed new attacks because signatures did not exist for these attacks and existing signatures did not generalize to variants of old attacks or to new and stealthy attacks. For example “ncftp” and “lsdomain” attacks were visible in tcpdump data, but missed because no rules existed to detect these attacks. Stealthy probes were missed because hard thresholds in rules were set to issue an alert only for more rapid probes, even though slow probes often provided as much information to attackers. Stealthy U2R attacks were missed by network-based systems because rules generated for clear versions of these attacks did not generalize to stealthy versions and because attacker actions were not easily visible in sniffing data. Many of the Windows NT attacks were missed due to lack of experience with Windows NT audit data and attacks. A detailed analysis of the Windows NT attacks [10] indicated that all but two of these attacks (ppmacro, framespoof) can be detected from the 1999 NT audit data using attack-specific signatures which generate far fewer than 10 false alarms per day.

Systems also missed attacks because particular protocols or services were not monitored. For example, some systems missed the “arppoisson” attack because the ARP protocol was not monitored. Some missed the “snmpget” attack because the SNMP service was not analyzed and some missed the “lsdomain” attack because the DNS service was not analyzed. Finally, some systems missed attacks because a protocol or TCP service was not analyzed to the required depth. For example, the “lsdomain” attack requires a system to monitor traffic to the DNS server and also detect when an “ls” command is successfully run on that server. The “selfping” command also will not be detected by a network-based intrusion detection system unless telnet sessions are extracted and analyzed to detect when a “ping” command is issued with specific arguments.

Some inside attacks were launched from the consoles of victims and did not generate network traffic. They were detected well only on the Solaris victim by systems that use BSM audit data. Other inside machine-to-machine attacks were detected as well using inside sniffer data as attacks initiated from outside machines. One anomaly detection system [7] provided good results. It analyzed system-call sequences extracted from BSM audit data and provided a high detection rate similar to that of the best signature-based systems for Solaris U2R attacks, as shown in the upper right of Table 5.

The forensic information provided in identification list files was generally accurate for attacks that were correctly detected. Table 6 shows results for four high-performance systems that provided all optional identification information. The first column in this table shows the system type. The second column shows the total number of attacks detected by each system (at the highest false-alarm rate) followed by a slash and the number of in-spec attacks that this system should have detected as specified in the system description. The first two expert systems both detected roughly 80 attacks each. They were combined systems that could have detected a maximum of roughly 170 in-spec attacks using both host-based and network-based input data. The third system used network sniffing data alone and thus had fewer in-spec attacks

(102) and the fourth system used only Solaris file-system information and thus had only 27 in-spec attacks. The remaining columns show the accuracy of the identification information provided for detected attacks. The third column shows the percentage of detected attacks where the attack category label was correct. The fourth column shows the percentage of detected attacks where the names of old attacks were correct. Participants were provided a list of names of old attacks, which was used to label attacks before the evaluation was run. Items in this column apply only to old attacks that were detected. The next column shows the percentage of detected attacks where 90% or more of the victim ports were identified and the final column shows the percentage of detected attacks where all the source IP addresses were correctly identified.

This table shows that the additional identification information provided was generally accurate for attacks that were correctly detected. For example, for the first expert system, the attack category and name is correct roughly 90% of the time and the victim ports and source IP addresses are correctly identified for more than 70% of the detected attacks. The upper three systems in Table 6 all used network sniffing data and provided good identification performance. The last forensic analysis system was a host-based system. Its good performance suggests that much of the identification information required can be obtained from a host-based analysis that does not rely on audit data.

All systems in Table 6 also provided attack start times as optional identification information. These times were computed by participating systems using off-line data with no constraints on look-ahead and thus they do not necessarily represent times that could be provided by real-time system im-

plementations. Start time accuracy was generally good for R2L and DoS attacks. The attack start time latencies were less than 15 s for more than 80% of these attacks. Start time accuracy was not as good, and differed across systems for probe and U2R attacks. Start times were provided for probe attacks by the first three systems in Table 6. The third system (DMine) correctly identified the start of all probes to within 15 s while the first two expert systems had start time latencies that were often many minutes delayed for slower probes that spanned long time intervals.

The first two expert systems and the last system in Table 6 provided start times for U2R attacks. These attacks were unique because many of them included multiple, separate, telnet interactions separated by long time intervals, while others were performed as part of long, single, telnet sessions containing many normal user commands. In attacks that included multiple telnet sessions, initial sessions were run at user privilege level to prepare for the attack. The actual attack, which provided root-level privilege on UNIX machines, was run only in following sessions. Results for the first two expert systems in Table 6 and for the last forensic analysis system differed dramatically for these U2R attacks. The first two systems detected the time instant where the attacker became root, while the forensic analysis system traced the beginning of the attack either to the beginning of the first session where attack setup actions occurred or to the beginning of the telnet session where the attack occurred. Start times for six of the eight U2R attacks detected by the forensic analysis system were within 15 s of true start times, while start times for more than 90% of the U2R attacks detected by the first two expert systems were delayed by more than a minute from the true attack times. These results suggest that the forensic analysis system accurately

Table 6  
Identification results for all attacks by four high-performance systems which provided optional identification information

	Attacks detected/ in-spec	Attack categories correct (%)	Attack names correct (old attacks, %)	Victim ports correct (%)	Source IP addresses correct (%)
Expert-1	85/169	91	88	73	80
Expert-2	81/173	74	53	51	69
DMine	41/102	100	88	61	90
Forensics	15/27	87	69	87	73



correlates information across multiple network sessions to arrive at accurate start times while the two expert systems use the time of the root-privilege elevation as a start time.

## 9. Discussion

The DARPA 1999 intrusion detection evaluation successfully evaluated 18 intrusion detection systems from 8 sites using more than 200 instances of 58 attack types embedded in three weeks of training data and two weeks of test data. Attacks were primarily launched against UNIX and Windows NT hosts. The best detection was provided by network-based systems for old probe and old DoS attacks and by host-based systems for Solaris U2R attacks launched either remotely or from the local console. A number of sites developed systems that detect known old attacks by searching for signatures in network sniffer data or Solaris BSM audit data using expert systems or rules. These systems detect old attacks well when they match known signatures, but miss many new UNIX attacks, Windows NT attacks and stealthy attacks. Promising capabilities were provided by Solaris host-based systems which detected console-based and remote-stealthy U2R attacks, anomaly detection systems which could detect some U2R and DoS attacks without requiring signatures and a host-based system that could detect Solaris U2R and R2L attacks without using audit information but by performing a forensic analysis of the Solaris file system.

Results of the 1999 evaluation should be interpreted within the context of the test bed, background traffic, attacks and scoring procedures used. The evaluation used a reasonable, but not exhaustive, set of attacks with a limited set of actions performed as part of each attack. It also used a simple network topology, a non-restrictive security policy, a limited number of victim machines and intrusion detection systems, stationary and low-volume background traffic, lenient scoring and extensive instrumentation to provide inputs to intrusion detection systems. One finding that should not be misinterpreted is that most systems had false-alarm rates which were low and well

below 10 false alarms per day. As noted above, these low rates may be caused by the use of relatively low volume background traffic with a time-varying, but relatively fixed, proportion of different traffic types. We currently plan to verify false-alarm rates using live network traffic and a small number of high-performing systems. Live-traffic measurements will also be made to update traffic statistics and traffic generators used in the test bed. Results obtained with the DARPA research systems used in the evaluation also may not generalize to more recent research systems or commercial systems. Performance with the 56 attack types used in the evaluation also may not be representative of performance with more recent attacks or with other attacks against different host machines, firewalls, routers or parts of the network infrastructure. Further evaluations are required to explore performance with commercial and other research intrusion detection systems, more complex network topologies, a wider range of attacks and varying mixtures and amounts of background traffic.

Comprehensive evaluations of DARPA research systems have now been performed in 1998 and 1999. These evaluations take time and effort on the part of the evaluators and the participants. They have provided benchmark measurements that do not now need to be repeated until system developers are able to implement many desired improvements. The current, planned, short-term focus in 2000 is to provide assistance to intrusion detection system developers to advance their systems and not to evaluate performance. System development can be expedited by providing descriptions and labeled examples of many new attacks, developing threat and attack models and carefully evaluating COTS systems to determine where to focus research efforts.

A number of approaches to improve capabilities of existing systems are suggested by the 1999 results. First, techniques should be developed to process Windows NT audit data and detect attacks by extending existing approaches from UNIX to Windows NT. Second, host-based systems should not rely exclusively on C2-level audit data such as Solaris BSM data or NT audit data. Instead, they should also examine information in the file system

and in commonly-used system logs. Systems that use file system information could be used on hosts such as Linux, where there currently is no C2-level auditing, and on any critical host where auditing is not turned on for fear of performance degradation. Third, systems should analyze a wider range of protocols and TCP services. For some protocols, information contained in packet headers alone is insufficient, but the content of network transmissions must be extracted to determine the purpose of important network interactions. Fourth, approaches that can detect new attacks, including anomaly detection, should be extended to more hosts and network traffic types. Fifth, systems should provide more forensic information to analysts and extend the optional attack identification information provided by many systems in 1999. This forensic analysis could simplify the task of verifying each alert, determining attacker actions and responding to an attack. It could also provide a valuable lasting record of attack-related events. Finally, other types of input features should be explored. These could be provided by new system auditing software, firewall or router audit logs, SNMP queries, software wrappers and application-specific auditing.

### Acknowledgements

This work was sponsored by the Department of Defense Advanced Research Projects Agency under Air Force Contract F19628-95-C-0002. Opinions, interpretations, conclusions and recommendations are those of the authors and not necessarily endorsed by the US Air Force. We would like to thank Sami Saydjari for supporting this effort. Many involved participants made this evaluation possible, including Dick Kemmerer, Giovanni Vigna, Mabri Tyson, Phil Porras, Anup Ghosh, R.C. Sekar and NingNing Wu. We would also like to thank Terry Champion and Steve Durst from AFRL for providing Linux kernel modifications that make one host simulate many IP addresses. Finally, we would like to thank others who contributed, including Marc Zissman, Rob Cunningham, Seth Webster, Kris Kendall, Raj Basu, Jesse Rabek and Simson Garfinkel.

### References

- [1] J. Allen, A. Christie, W. Fithen, J. McHugh, J. Pickel, E. Stoner, State of the practice of intrusion detection technologies, Carnegie Mellon University/Software Engineering Institute Technical Report CMU/SEI-99-TR-028, January 2000.
- [2] E.G. Amoroso, *Intrusion Detection: An Introduction to Internet Surveillance, Correlation, Trace Back, Traps, and Response*, Intrusion.Net Books, 1999.
- [3] K. Das, The development of stealthy attacks to evaluate intrusion detection systems, S.M. Thesis, MIT Department of Electrical Engineering and Computer Science, June 2000.
- [4] H. Debar, M. Dacier, A. Wespi, S. Lampart, An experimental workbench for intrusion detection systems, Research Report RZ 2998 (#93044), IBM Research Division, Zurich Research Laboratory, 8803 Ruschlikon, Switzerland, 9 March 1999, <http://www.zurich.ibm.com/Technology/Security/extern/gsal/docs/index.html>.
- [5] R. Durst, T. Champion, B. Witten, E. Miller, L. Spagnuolo, Testing and evaluating computer intrusion detection systems, *Communications of the ACM* 42 (1999) 53–61.
- [6] C. Elkan, Results of the KDD'99 Classifier Learning Contest, Sponsored by the International Conference on Knowledge Discovery in Databases, September 1999, <http://www-cse.ucsd.edu/users/elkan/clresults.html>.
- [7] A.K. Ghosh, A. Schwartzbard, A study in using neural networks for anomaly and misuse detection, in: *Proceedings of the USENIX Security Symposium*, 23–26 August 1999, Washington, DC, <http://www.rstcorp.com/~anup/>.
- [8] S. Jajodia, D. Barbara, B. Speegle, N. Wu, Audit Data Analysis and Mining (ADAM), project described in <http://www.isse.gmu.edu/~dbarbara/adam.html>, April 2000.
- [9] K. Kendall, A database of computer attacks for the evaluation of intrusion detection systems, S.M. Thesis, MIT Department of Electrical Engineering and Computer Science, June 1999.
- [10] J. Korba, Windows NT attacks for the evaluation of intrusion detection systems, S.M. Thesis, MIT Department of Electrical Engineering and Computer Science, June 2000.
- [11] R.P. Lippmann, David J. Fried, Isaac Graf, Joshua W. Haines, Kristopher R. Kendall, David McClung, Dan Weber, Seth E. Webster, Dan Wyschogrod, Robert K. Cunningham, Marc A. Zissman, Evaluating intrusion detection systems: the 1998 DARPA off-line intrusion detection evaluation, in: *Proceedings of the 2000 DARPA Information Survivability Conference and Exposition (DISCEX)*, vol. 2, IEEE Press, New York, January 2000.
- [12] R.P. Lippmann, R.K. Cunningham, Guide to creating stealthy attacks for the 1999 DARPA off-line intrusion detection evaluation, MIT Lincoln Laboratory Project Report IDDE-1, June 1999.
- [13] P. Neumann, P. Porras, Experience with EMERALD to DATE, in: *Proceedings of the First USENIX Workshop on Intrusion Detection and Network Monitoring*, Santa

- Clara, CA, April 1999, pp. 73–80, <http://www.sdl.sri.com/emerald/index.html>.
- [14] S. Northcutt, *Network Intrusion Detection; An Analysis Handbook*, New Riders, Indianapolis, 1999.
  - [15] T.H. Ptacek, T.N. Newsham, Insertion, Evasion, and Denial of Service: Eluding Network Intrusion Detection, Secure Networks, Inc. Report, January 1998.
  - [16] N. Puketza, K. Zhang, M. Chung, B. Mukherjee, R.A. Olsson, A methodology for testing intrusion detection systems, *IEEE Transactions on Software Engineering* 22 (1996) 719–729.
  - [17] N. Puketza, M. Chung, R.A. Olsson, B. Mukherjee, A software platform for testing intrusion detection systems, *IEEE Software* (September/October 1997) 43–51.
  - [18] A. Schwartzbard, A.K. Ghosh, A study in the feasibility of performing host-based anomaly detection on Windows NT, in: *Proceedings of the Second Recent Advances in Intrusion Detection (RAID 1999) Workshop*, West Lafayette, IN, 7–9 September 1999.
  - [19] R. Sekar, P. Uppuluri, Synthesizing fast intrusion prevention/detection systems from high-level specifications, in: *Proceedings of the Eighth Usenix Security Symposium*, Washington, DC, August 1999, <http://rcs-sgi.cs.iastate.edu/sekar/abs/usenixsec99.htm>.
  - [20] G. Shipley, Intrusion detection, take two, *Network Computing*, 15 November 1999, <http://www.nwc.com/1023/10231.html>.
  - [21] D. Song, G. Shaffer, M. Undy, Nidsbench – A network intrusion detection system test suite, in: *The Second International Workshop on Recent Advances in Intrusion Detection (RAID)*, September 1999, <http://www.anzen.com/research/nidsbench/nidsbench-slides/nidsbench-slides.html>.
  - [22] M. Tyson, P. Berry, N. Williams, D. Moran, D. Blei, DERBI: Diagnosis, Explanation and Recovery from Computer Break-Ins, project described in <http://www.ai.sri.com/~derbi/>, April 2000.
  - [23] G. Vigna, S.T. Eckmann, R.A. Kemmerer, The STAT tool suite, in: *Proceedings of the 2000 DARPA Information Survivability Conference and Exposition (DISCEX)*, IEEE Press, New York, January 2000.
  - [24] G. Vigna, R. Kemmerer, NetSTAT: a network-based intrusion detection system, *Journal of Computer Security* 7(1) (1999).

**Richard P. Lippmann** received a B.S. degree in Electrical Engineering from the Polytechnic Institute of Brooklyn in 1970 and a Ph.D. degree in Electrical Engineering from the Massachusetts Institute of Technology in 1978. From 1978 to 1981, he was Director of the Communications Engineering Laboratory of the Boys Town Institute for Communication Disorders in Children, Omaha, NE, working on speech perception, speech training aids for deaf children, sound alerting aids for the deaf and signal processing for hearing aids. In 1981, he joined the Massachusetts Institute of Technology, Lincoln Laboratory, and is currently a Senior Staff Member in the Information Systems Technology Group. Recent research interests include computer intrusion detection, speech recognition by humans and machines, developing improved neural network and statistical pattern classifiers, medical risk assessment, the development of public-domain software for pattern classification and the application of neural networks and statistics to problems in computer intrusion detection. He has supervised numerous MIT student theses in these areas. Dr. Lippmann is currently a Distinguished Lecturer for the IEEE Signal Processing Society.